

Pseudo-random Number Generation

Random Numbers

- All the code we've written is deterministic - it does the same thing every time.
- Some things we might want to simulate are apparently not so - nuclear decay, "noise" on signals, ages of a person you bump into.
- Sometimes we also just want to generate "different values", to explore the phase space of a problem.
- So, how do we generate "random variation"?

Pseudo Random Numbers

- Most of the time, we don't really care about having truly "random" sequences.
- What we want are sequences which have statistical properties of a random sequence:
 - each value is hard to predict from the ones before it (**uncorrelated**)
 - the probability of any value is the same (**uniform**)
- We can imagine a *pseudorandom* number generator, which produces such a sequence, but deterministically, from some kind of internal state.

A (bad) example PRNG

0**54**1

$$\mathbf{54}^2 = 2916$$

2**91**6

$$\mathbf{91}^2 = 8281$$

8**28**1

$$\mathbf{28}^2 = 0784$$

0**78**4

$$\mathbf{78}^2 = 6084$$

6**08**4

$$\mathbf{08}^2 = 0064$$

0**06**4

$$\mathbf{06}^2 = \dots$$

...

- John Von Neumann wrote this PRNG in 1949.
- We start with a "seed" value: here 0541.
- Each value in the sequence is the square of the middle digits of the value before it.

PRNGs - flaws

0000	0001	0002	9008		1453		1243
0000	0000	0000	0000		2025		0576
0000	0000	0000	0000		0004		3249
0000	0000	0000	0000	...	0000	...	0576
0000	0000	0000	0000		0000		3249
0000	0000	0000	0000		0000		0576
...

- Badly designed PRNGs can have unexpected behaviour for some seed values.
- Here, starting with any sequence containing 00, 45, 24 (or 57) generates a very predictable ("non-random") sequence!

PRNGs - flaws

- Even much more sophisticated PRNGs, which don't have obvious flaws, can have more subtle issues.
- For example: IBM's **RANDU** PRNG, which was widely used in the 1960s.
- Although the output "looked random", it actually had some extremely strong correlations, which made the output untrustworthy.
- Many millions of hours of simulations had to be redone once this was discovered!
- (Plus, obviously, no PRNG can ever be *truly* uncorrelated, they just approximate this state.)

PRNGs - advantages

- Because a PRNG generates a sequence from an internal state, which is set by a seed value...
- ...the same seed will always give the same sequence.
- So, we can directly compare the result of changes to code, by running it with the same seed as before.

PRNGs - advantages

- The alternative to PRNGs is to try to "collect" randomness from the real world.
- "Hardware Random Number Generators" exist, but must gather truly random sources - radioactive decay, or shot noise in electronic circuits - which limits the rate at which they can work.
- PRNGs are usually faster than HRNGs; and you can always rely on them being available!

PRNGs in C

- The C Standard Library provides some functions for generation of pseudorandom sequences.
- These functions are prototyped in `stdlib.h`, so you need to `#include` it to use them.
- `void srand(unsigned int seed)`
 - set seed for PRNG to provided seed value (default is `0`)
- `unsigned int rand(void)`
 - get next value from PRNG (integer in range `0` to `RAND_MAX`)
- `RAND_MAX`
 - `#defined` in `stdlib.h`: the largest value that `rand()` will ever return.

Example - continuous uniform variates


```
double zero_to_one(void) {  
    return rand() / (double) RAND_MAX;  
}
```

`rand()` will only generate values from
0 to **RAND_MAX**

If we re-scale those values by dividing by **RAND_MAX**,
then we will clearly get a number between **0** and **1**.
(Remember to cast to **double**, to avoid integer division)

Example (approximation of Gaussian / Normally distributed variates)

- Sometimes we want a *non-uniform* distribution.
- A common distribution we might want is the Gaussian distribution (the Normal distribution with $\mu = 0$, $\sigma = 1$).



There are efficient, exact, methods to produce arbitrary probability distributions, but these are not in scope for this course.

Example (approximation of Gaussian / Normally distributed variates)

- The **Central Limit Theorem** states that the distribution of *sums of a large number of samples from any distribution* itself approximates the Gaussian distribution, as the number of samples summed over tends to infinity.
- So, if we average over a large number of samples from a uniform distribution - for example, calls to **rand()**, and return the result...
- ...the resulting sequence will be an approximation to the distribution we want.

Example (approximation of Gaussian / Normally distributed variates)

```
#include <math.h>

//number of uniform deviates to add for our approximation. Larger is better, but slower
#define GAUSS_PREC 6

double approx_gauss(void) {

    double variate=0.0;

    for(int i=0;i<GAUSS_PREC;i++) {

        variate += rand() / (double) RAND_MAX;

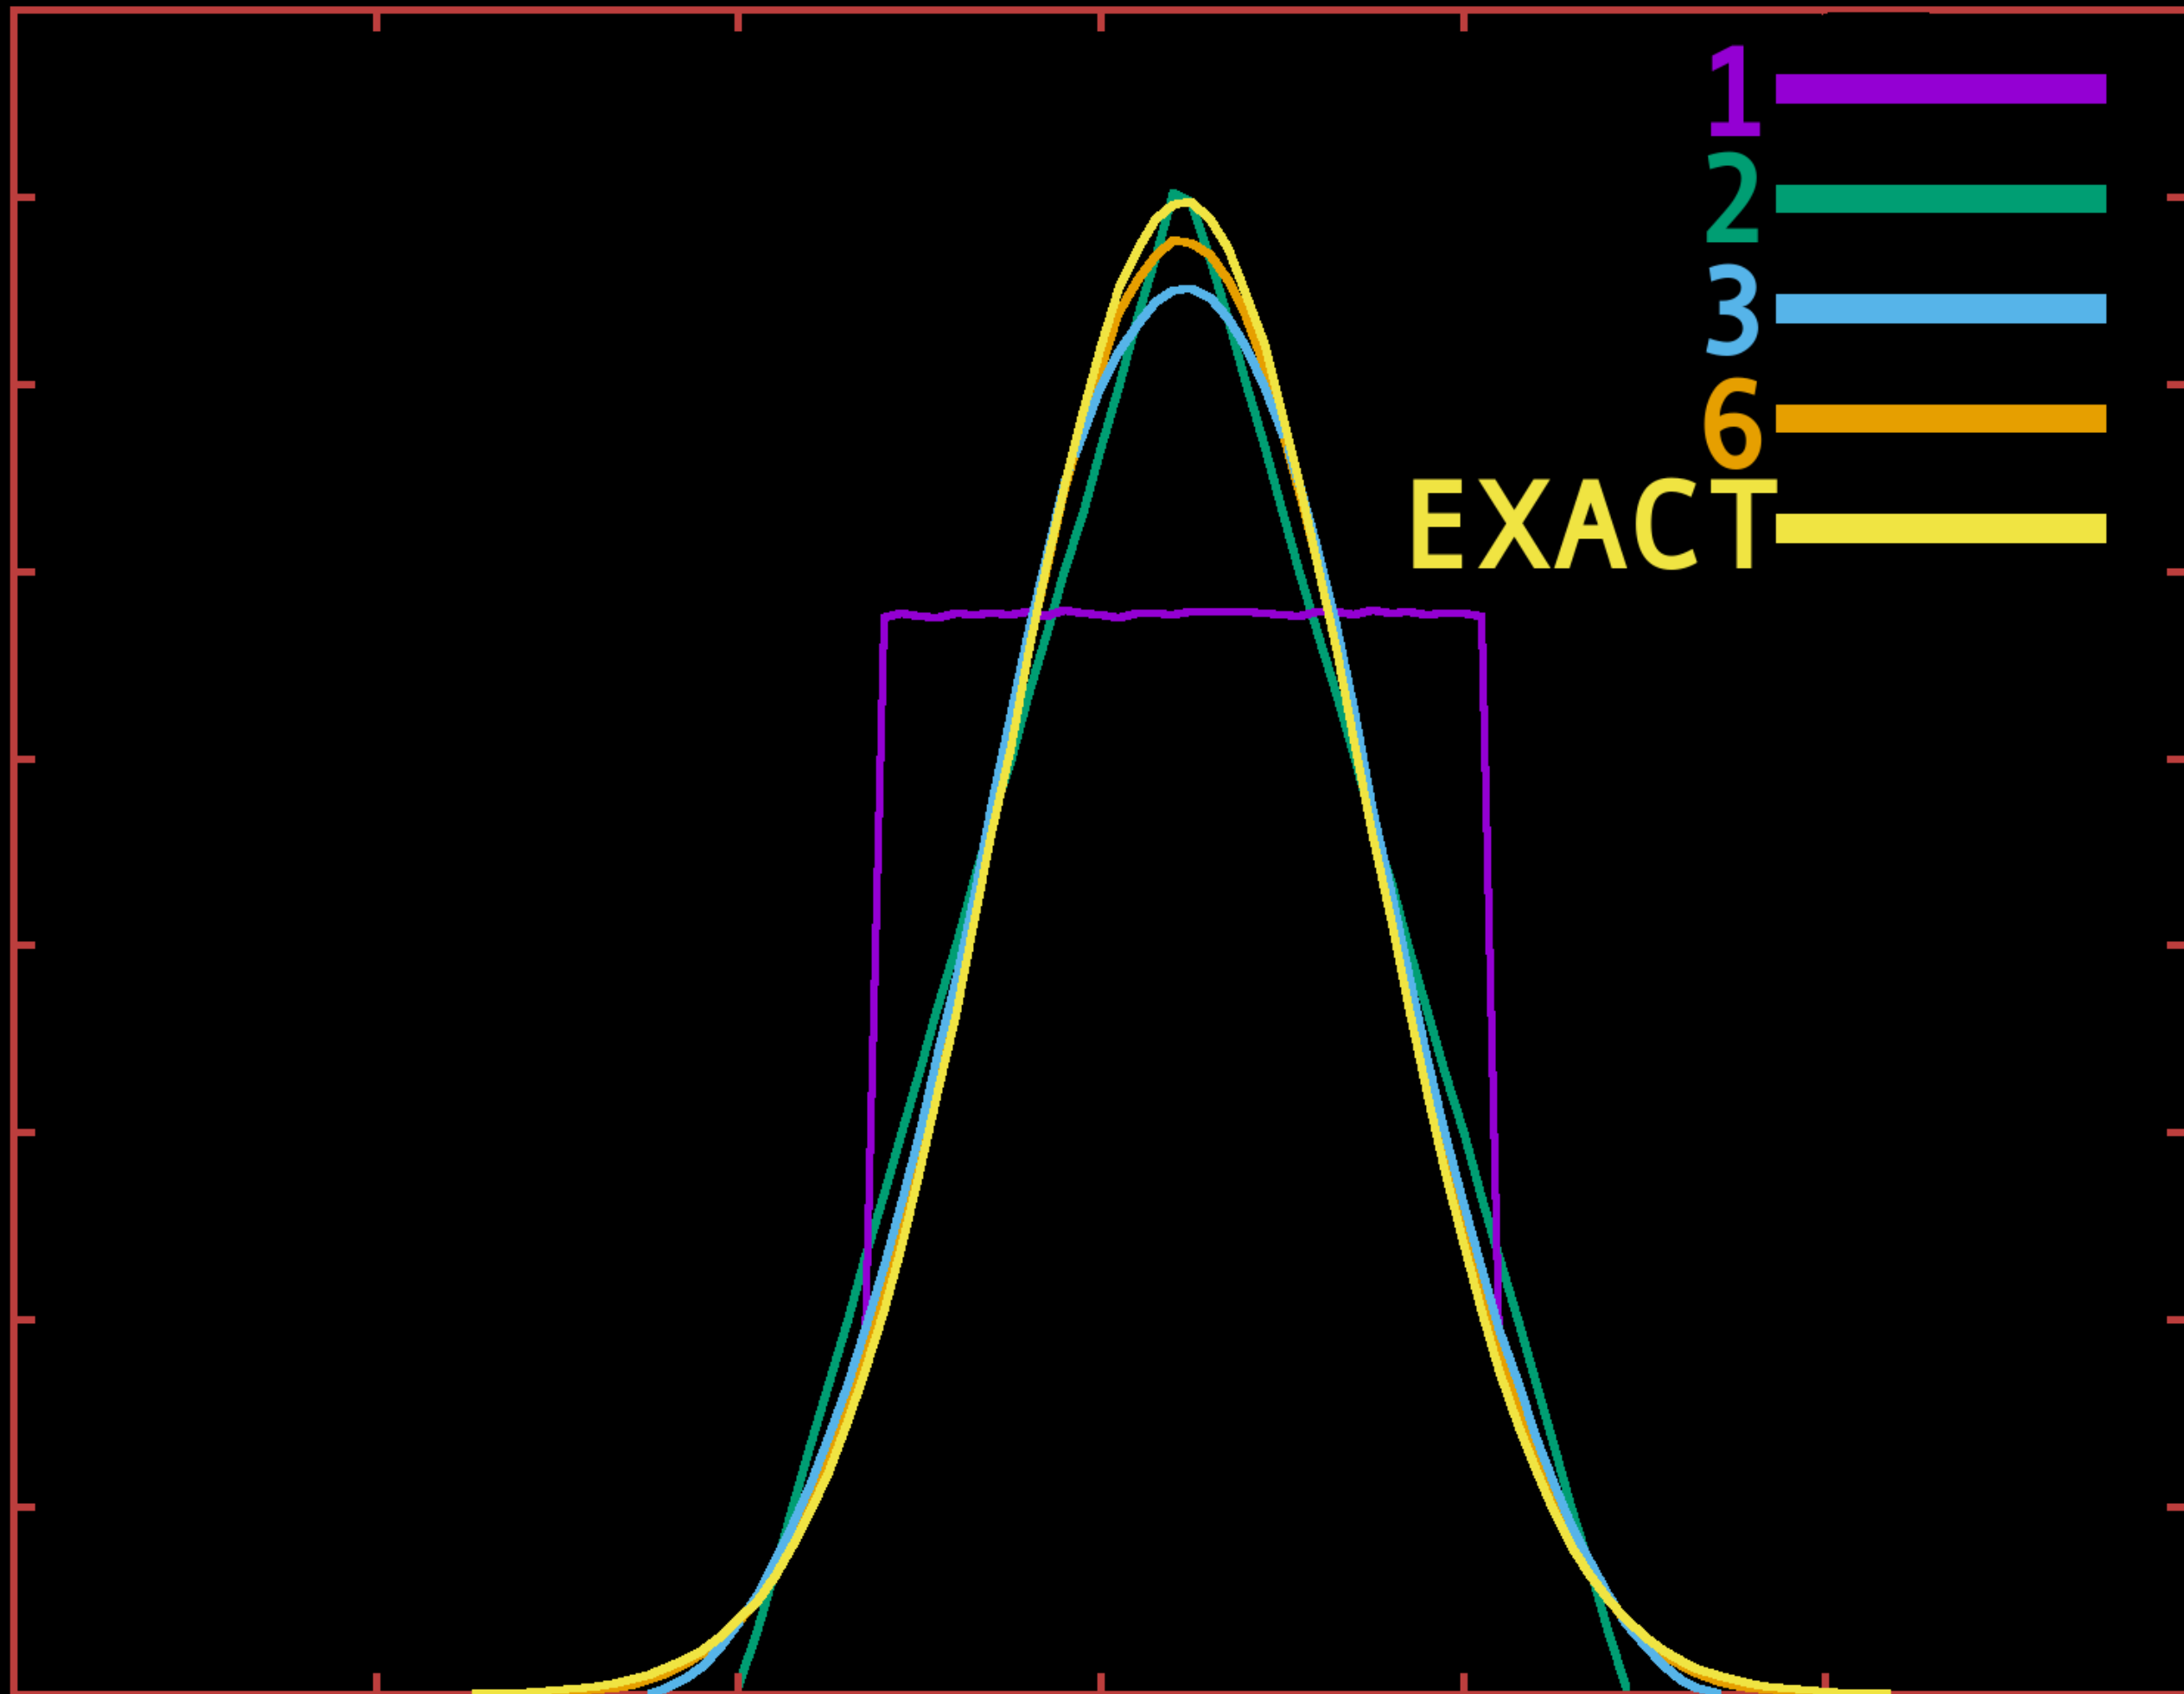
    }

    //rescale variance and mean appropriately for Gaussian
    // distribution is still Normal without this, just centred on GAUSS_PREC/2
    // with variance GAUSS_PREC / 12.0

    return (variate - 0.5*GAUSS_PREC) / sqrt(GAUSS_PREC/12.0);

}
```

This method takes GAUSS_PREC uniform variates to make 1 approximate Gaussian variate. More efficient, exact - such as the Box-Müller and Marsaglia's Polar - methods exist, but are not in scope for this course.



Example - seeding rand from system time.

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
```

time.h header for time related functions in libc

```
int main(void) {
    /* time(TIME) returns the current time as seconds
       since the "Unix Epoch" (midnight 1 Jan 1970).
       TIME can be NULL (or a pointer to a place to store
       the result, instead of returning it) */
```

```
    srand(time(NULL));
```

time(NULL) returns a number which will be different every second, "good enough" for the seed to be different on each execution

```
    //print 100 pseudo random numbers in sequence
    for(int i=0; i<100; i++) {
        printf("%d\n", rand() );
    }
```

```
    return 0;
}
```

Limitations of rand()

- `rand()` is convenient, being in the C Standard Library, but it has limitations.
- As `srand()` takes an unsigned int as a seed, there are only `UINT_MAX` possible sequences that `rand` can produce (one per seed).
- The C Standard doesn't require `rand` to have a particularly long sequence length (for portability). In some versions of the library, `rand` can have quite short repeat lengths, or fail some statistical tests.

Beyond rand()...

- For serious work requiring pseudorandom sequences with good statistical properties, you should use a dedicated PRNG from a library designed for this.
- Mersenne Twister and WELL, for example, are two widely used and very good PRNGs which you should consider.
- (For cryptography, even these are not sufficient - consider special cryptographically strong PRNGs, or actually taking random numbers from outside machine.)

References

- Mersenne Twister
- <http://www.math.sci.hiroshima-u.ac.jp/~m-mat/MT/emt.html>
- WELL
- <http://www.iro.umontreal.ca/~panneton/WELLRNG.html>
- Example use for PRNGs:
- <http://moodle2.gla.ac.uk/mod/resource/view.php?id=149996>

An aside on pointers.

- You can have pointers to anything.
- This includes *functions*.
 - functions exist in memory.
 - And they have a type (their signature).

```
int myfunc (int y, double z); // a function prototype
```

```
int (*myfunc_ptr) (int, double);
```

```
//myfunc_ptr is a pointer to functions with myfunc's signature.
```

```
// it *CANNOT* safely point to functions with a different signature.
```

```
myfunc_ptr = myfunc; //point myfunc_ptr at myfunc
```

```
myfunc_ptr(1,2.0); //the same as myfunc(1,2.0)
```

Using function pointers to make a generic test

need the () around the * so compiler knows we mean "pointer to function which returns int" and not "function which returns pointer to int"

```
double mean( int (*f) (void) ) {  
    //f here is a "pointer to a function which takes void  
    //                                     and returns double"  
    double accum = 0;  
    for (int i=0; i<50; i++) {  
        accum += f(); //you can call the function pointer with () as normal  
    }  
    return accum/50;  
}
```

a function name without the () is a pointer to it
(just like the array name without the [] is a pointer)

```
double mu = mean(rand); //return mean of 50 invocations of rand()
```